

```
from sklearn.datasets import load_iris
iris = load_iris()
```

```
x = iris.data
y = iris.target
```

```
import pandas as pd
df = pd.DataFrame(x, columns = iris.feature_names)
```

```
df['target'] = y
print(df.head())
print()
print(df.describe())
```

```
↩
  sepal length (cm)  sepal width (cm)  petal length (cm)  petal width (cm)  \
0                5.1                3.5                1.4                0.2
1                4.9                3.0                1.4                0.2
2                4.7                3.2                1.3                0.2
3                4.6                3.1                1.5                0.2
4                5.0                3.6                1.4                0.2
```

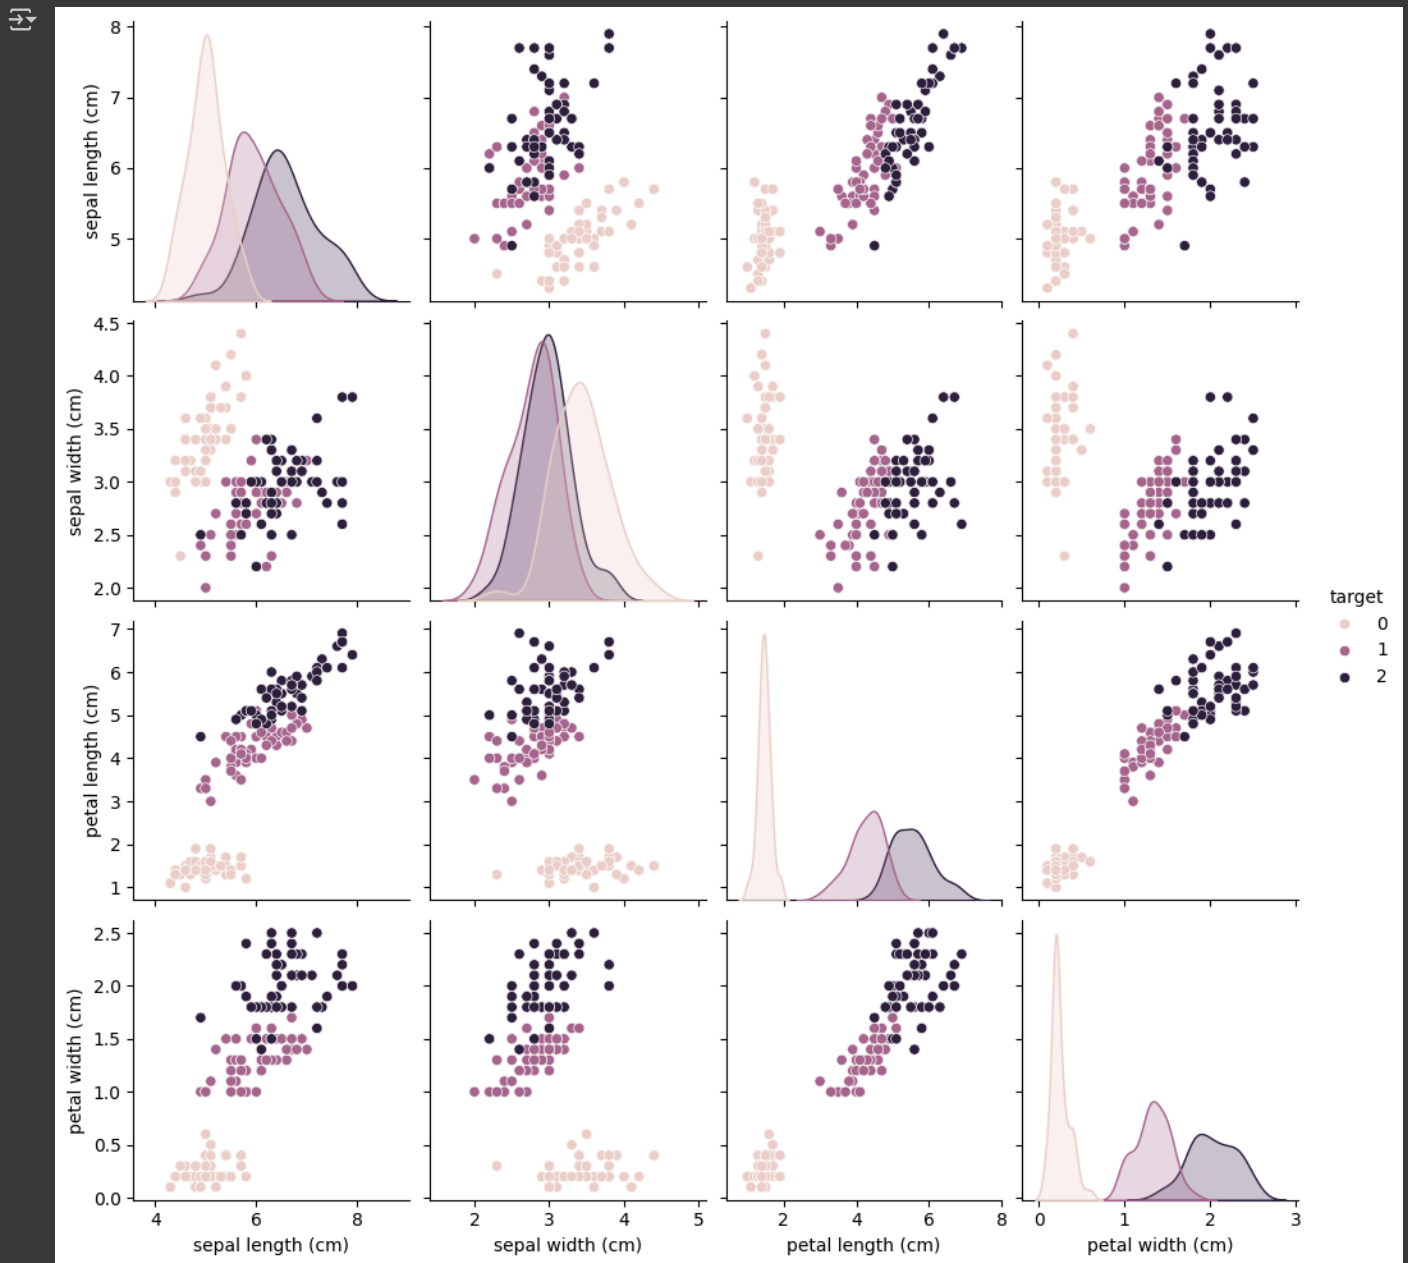
```
target
0      0
1      0
2      0
3      0
4      0
```

```
      sepal length (cm)  sepal width (cm)  petal length (cm)  \
count      150.000000      150.000000      150.000000
mean         5.843333         3.057333         3.758000
std          0.828066         0.435866         1.765298
min          4.300000         2.000000         1.000000
25%          5.100000         2.800000         1.600000
50%          5.800000         3.000000         3.300000
75%          6.400000         3.300000         5.100000
max          7.900000         4.400000         6.900000
```

```
      petal width (cm)  target
count      150.000000  150.000000
mean         1.199333     1.000000
std          0.762238     0.819232
min          0.100000     0.000000
25%          0.300000     0.000000
50%          1.300000     1.000000
75%          1.800000     2.000000
max          2.500000     2.000000
```

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
sns.pairplot(df, hue="target")
plt.show()
```



```
from sklearn.neural_network import MLPClassifier
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=
scaler = StandardScaler()
x_train_scaled = scaler.fit_transform(x_train)
x_test_scaled = scaler.transform(x_test)

clf = MLPClassifier(hidden_layer_sizes=(5, 3), activation='relu', sc
clf.fit(x_train_scaled, y_train)
```

MLPClassifier

MLPClassifier(hidden_layer_sizes=(5, 3), max_iter=3000, random_state=42)

```
y_pred = clf.predict(x_test_scaled)
```

```
print(y_pred)
```

```
[ 1  0  0  0  1  1  0  2  0  0  1  1  2  2  0  2  2  1  0  2  2  1  2  1  0  1  0  0  1  2]
```

```
from sklearn.metrics import confusion_matrix, accuracy_score, classi
```

```
print(confusion_matrix(y_test, y_pred))
```

```
print()
```

```
print(accuracy_score(y_test, y_pred))
```

```
print()
```

```
print(classification_report(y_test, y_pred))
```

```
[[11  0  0]
 [ 0 10  2]
 [ 0  0  7]]
```

```
0.9333333333333333
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	11
1	1.00	0.83	0.91	12
2	0.78	1.00	0.88	7
accuracy			0.93	30
macro avg	0.93	0.94	0.93	30
weighted avg	0.95	0.93	0.93	30

```
plt.figure(figsize=(6,4))
```


```
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, cmap='Blue
```

```
plt.xlabel('Predicted')
```

```
plt.ylabel('Actual')
```

```
plt.title('Confusion Matrix')
```

```
plt.show()
```



```
from sklearn.model_selection import cross_val_score
```

```
scores = cross_val_score(clf, scaler.transform(x), y, cv=5)
```

```
print("Cross-validation scores:", scores)
```

```
print("Mean CV score:", scores.mean())
```

↗ Cross-validation scores: [0.96666667 1. 0.9 0.93333333 0.96666667]
Mean CV score: 0.9533333333333334

```
from sklearn.model_selection import GridSearchCV
from sklearn.neural_network import MLPClassifier

param_grid = {
    'hidden_layer_sizes': [(5, 5), (10, 3), (10, 5), (15, 3)],
    'activation': ['relu', 'logistic', 'tanh'],
    'solver': ['adam', 'sgd'],
    'learning_rate_init': [0.0005, 0.001, 0.01, 0.1],
    'max_iter': [1000, 2000, 3000]
}

clf = MLPClassifier(random_state=42)

grid_search = GridSearchCV(clf, param_grid, cv=5, n_jobs=-1, verbose=1)
grid_search.fit(scaler.transform(x), y)

print("Best Parameters:", grid_search.best_params_)
print("Best Cross-Validation Score:", grid_search.best_score_)
```

↗ Fitting 5 folds for each of 288 candidates, totalling 1440 fits
Best Parameters: {'activation': 'relu', 'hidden_layer_sizes': (5, 5), 'learning_rate_init': 0.0005, 'max_iter': 2000, 'solver': 'adam'}
Best Cross-Validation Score: 0.9733333333333334

Use best Parameters from GridSearchCV

```
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2)

scaler = StandardScaler()
x_train_scaled = scaler.fit_transform(x_train)
x_test_scaled = scaler.transform(x_test)

clf = MLPClassifier(hidden_layer_sizes=(5, 5), activation='relu', solver='adam',
                    learning_rate_init=0.0005, max_iter=2000, random_state=42)
clf.fit(x_train_scaled, y_train)
```

↗

MLPClassifier

MLPClassifier(hidden_layer_sizes=(5, 5), learning_rate_init=0.0005, max_iter=2000, random_state=42)

```
y_pred = clf.predict(x_test_scaled)

print(y_pred)
```

↗ [1 0 0 0 1 1 0 2 0 0 1 1 2 2 0 2 2 1 0 2 2 1 2 1 0 1 0 0 1 2]

```

print(confusion_matrix(y_test, y_pred))
print()
print(accuracy_score(y_test, y_pred))
print()
print(classification_report(y_test, y_pred))

```

```

[[11  0  0]
 [ 0 10  2]
 [ 0  0  7]]

```

```
0.9333333333333333
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	11
1	1.00	0.83	0.91	12
2	0.78	1.00	0.88	7
accuracy			0.93	30
macro avg	0.93	0.94	0.93	30
weighted avg	0.95	0.93	0.93	30

```

from sklearn.model_selection import cross_val_score
scores = cross_val_score(clf, scaler.transform(x), y, cv=5)

```